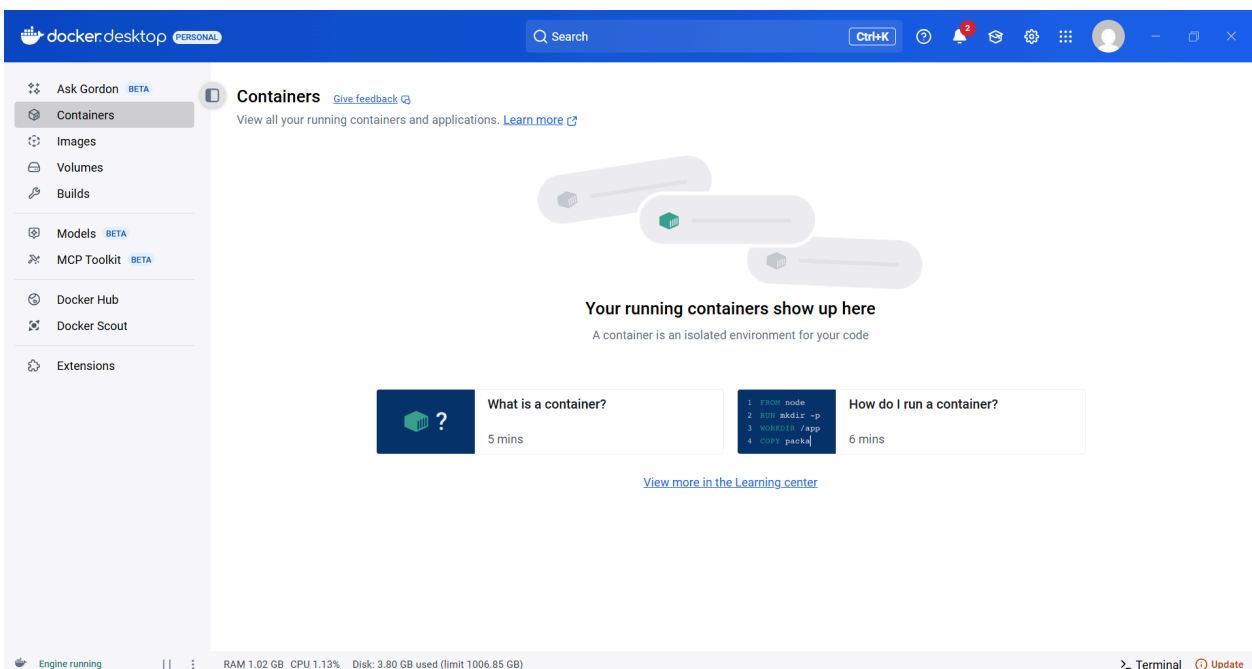


Airflow (2.4.2) on Windows + Docker - Manual

1) Prerequisites

1. Install **Docker Desktop for Windows** (with WSL2 backend enabled).

<https://www.docker.com/products/docker-desktop/>



2. Open **PowerShell** (as a normal user is fine).

3. Create a working folder:

```
mkdir airflow  
cd airflow
```

This will be the working directory for Airflow.

2) Get the Docker Compose file

- Use the **CeleryExecutor** example for Airflow **2.4.2** (Postgres + Redis).

- **Important edits we made** to avoid errors:
 - Removed the **airflow-apiserver** service (Airflow 2.4.2 doesn't have the **api-server** command).
 - Removed any **depends_on: airflow-apiserver** (e.g., in **airflow-worker**).
 - Fixed the **duplicate healthcheck key** (YAML cannot define the same key twice in one block).

Ensured the **webserver** has:

- **ports:**
 - "8080:8080"

YAML file :

```
# Licensed to the Apache Software Foundation (ASF) under one
# or more contributor license agreements. See the NOTICE file
# distributed with this work for additional information
# regarding copyright ownership. The ASF licenses this file
# to you under the Apache License, Version 2.0 (the
# "License"); you may not use this file except in compliance
# with the License. You may obtain a copy of the License at
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing,
# software distributed under the License is distributed on an
# "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY
# KIND, either express or implied. See the License for the
# specific language governing permissions and limitations
# under the License.
#

# Basic Airflow cluster configuration for CeleryExecutor with Redis and PostgreSQL.
#
# WARNING: This configuration is for local development. Do not use it in a production deployment.
#
# This configuration supports basic configuration using environment variables or an .env file
# The following variables are supported:
#
# AIRFLOW_IMAGE_NAME           - Docker image name used to run Airflow.
#                               Default: apache/airflow:3.0.4
# AIRFLOW_UID                  - User ID in Airflow containers
#                               Default: 50000
# AIRFLOW_PROJ_DIR              - Base path to which all the files will be volumed.
#                               Default: .
# Those configurations are useful mostly in case of standalone testing/running Airflow in test/try-out mode
#
```

```

# _AIRFLOW_WWW_USER_USERNAME - Username for the administrator account (if requested).
#                               Default: airflow
# _AIRFLOW_WWW_USER_PASSWORD - Password for the administrator account (if requested).
#                               Default: airflow
# _PIP_ADDITIONAL_REQUIREMENTS - Additional PIP requirements to add when starting all containers.
#                               Use this option ONLY for quick checks. Installing requirements at container
#                               startup is done EVERY TIME the service is started.
#                               A better way is to build a custom image or extend the official image
#                               as described in https://airflow.apache.org/docs/docker-stack/build.html.
#                               Default: ''
#
# Feel free to modify this file to suit your needs.
---
x-airflow-common:
  &airflow-common
  # In order to add custom dependencies or upgrade provider distributions you can use your extended image.
  # Comment the image line, place your Dockerfile in the directory where you placed the docker-compose.yaml
  # and uncomment the "build" line below, Then run `docker-compose build` to build the images.
  image: ${AIRFLOW_IMAGE_NAME:-apache/airflow:3.0.4}
  # build: .
  environment:
    &airflow-common-env
    AIRFLOW__CORE__EXECUTOR: CeleryExecutor
    AIRFLOW__CORE__AUTH_MANAGER: airflow.providers.fab.auth_manager.fab_auth_manager.FabAuthManager
    AIRFLOW__DATABASE__SQL_ALCHEMY_CONN: postgresql+psycopg2://airflow:airflow@postgres/airflow
    AIRFLOW__CELERY__RESULT_BACKEND: db+postgresql://airflow:airflow@postgres/airflow
    AIRFLOW__CELERY__BROKER_URL: redis://:@redis:6379/0
    AIRFLOW__CORE__FERNET_KEY: ''
    AIRFLOW__CORE__DAGS_ARE_PAUSED_AT_CREATION: 'true'
    AIRFLOW__CORE__LOAD_EXAMPLES: 'true'
    AIRFLOW__CORE__EXECUTION_API_SERVER_URL: 'http://airflow-apiserver:8080/execution/'
  # yamllint disable rule:line-length

```

```

# Use simple http server on scheduler for health checks
# See https://airflow.apache.org/docs/apache-airflow/stable/administration-and-deployment/logging-monitoring/check-health.html#scheduler-health
# yamllint enable rule:line-length
AIRFLOW__SCHEDULER__ENABLE_HEALTH_CHECK: 'true'
# WARNING: Use _PIP_ADDITIONAL_REQUIREMENTS option ONLY for a quick checks
# for other purpose (development, test and especially production usage) build/extend Airflow image.
_PIP_ADDITIONAL_REQUIREMENTS: ${_PIP_ADDITIONAL_REQUIREMENTS:-}
# The following line can be used to set a custom config file, stored in the local config folder
AIRFLOW_CONFIG: '/opt/airflow/config/airflow.cfg'
volumes:
  - ${AIRFLOW_PROJ_DIR:-.}/dags:/opt/airflow/dags
  - ${AIRFLOW_PROJ_DIR:-.}/logs:/opt/airflow/logs
  - ${AIRFLOW_PROJ_DIR:-.}/config:/opt/airflow/config
  - ${AIRFLOW_PROJ_DIR:-.}/plugins:/opt/airflow/plugins
user: "${AIRFLOW_UID:-50000}:0"
depends_on:
  &airflow-common-depends-on
  redis:
    condition: service_healthy
  postgres:
    condition: service_healthy

services:
  postgres:
    image: postgres:13
    environment:
      POSTGRES_USER: airflow
      POSTGRES_PASSWORD: airflow
      POSTGRES_DB: airflow
    volumes:

```

```

- postgres-db-volume:/var/lib/postgresql/data
healthcheck:
  test: ["CMD", "pg_isready", "-U", "airflow"]
  interval: 10s
  retries: 5
  start_period: 5s
  restart: always

redis:
  # Redis is limited to 7.2-bookworm due to licencing change
  # https://redis.io/blog/redis-adopts-dual-source-available-licensing/
  image: redis:7.2-bookworm
  expose:
    - 6379
  healthcheck:
    test: ["CMD", "redis-cli", "ping"]
    interval: 10s
    timeout: 30s
    retries: 50
    start_period: 30s
    restart: always

airflow-webserver:
  <<: *airflow-common
  command: webserver
  ports:
    - "8080:8080"
  healthcheck:
    test: ["CMD", "curl", "--fail", "http://localhost:8080/health"]

```

```

    interval: 30s
    timeout: 10s
    retries: 5
    start_period: 30s
  restart: always
  depends_on:
    <<: *airflow-common-depends-on
    airflow-init:
      condition: service_completed_successfully

airflow-scheduler:
  <<: *airflow-common
  command: scheduler
  healthcheck:
    test: ["CMD", "curl", "--fail", "http://localhost:8974/health"]
    interval: 30s
    timeout: 10s
    retries: 5
    start_period: 30s
  restart: always
  depends_on:
    <<: *airflow-common-depends-on
    airflow-init:
      condition: service_completed_successfully

airflow-dag-processor:
  <<: *airflow-common
  command: dag-processor

```

```

healthcheck:
  test: ["CMD-SHELL", 'airflow jobs check --job-type DagProcessorJob --hostname "${HOSTNAME}"]
  interval: 30s
  timeout: 10s
  retries: 5
  start_period: 30s
restart: always
depends_on:
  <<: *airflow-common-depends-on
  airflow-init:
    condition: service_completed_successfully

airflow-worker:
  <<: *airflow-common
  command: celery worker
  healthcheck:
    # yamllint disable rule:line-length
    test:
      - "CMD-SHELL"
      - 'celery --app airflow.providers.celery.executors.celery_executor.app inspect ping -d "celery@${HOSTNAME}" || celery --app airflow.executo
    interval: 30s
    timeout: 10s
    retries: 5
    start_period: 30s
  environment:
    <<: *airflow-common-env
    # Required to handle warm shutdown of the celery workers properly
    # See https://airflow.apache.org/docs/docker-stack/entrypoint.html#signal-propagation
    DUMB_INIT_SETSID: "0"
  restart: always

```

```

  depends_on:
    <<: *airflow-common-depends-on
    airflow-webserver:
      condition: service_healthy
    airflow-init:
      condition: service_completed_successfully

airflow-triggerer:
  <<: *airflow-common
  command: triggerer
  healthcheck:
    test: ["CMD-SHELL", 'airflow jobs check --job-type TriggererJob --hostname "${HOSTNAME}"]
    interval: 30s
    timeout: 10s
    retries: 5
    start_period: 30s
  restart: always
  depends_on:
    <<: *airflow-common-depends-on
    airflow-init:
      condition: service_completed_successfully

airflow-init:
  <<: *airflow-common
  entrypoint: /bin/bash
  # yamllint disable rule:line-length
  command:
    - -c
    - |
      if [[ -z "${AIRFLOW_UID}" ]]; then

```

```

if [[ -z "${AIRFLOW_UID}" ]]; then
    echo
    echo -e "\033[1;33mWARNING!!!: AIRFLOW_UID not set!\e[0m"
    echo "If you are on Linux, you SHOULD follow the instructions below to set "
    echo "AIRFLOW_UID environment variable, otherwise files will be owned by root."
    echo "For other operating systems you can get rid of the warning with manually created .env file:"
    echo "    See: https://airflow.apache.org/docs/apache-airflow/stable/howto/docker-compose/index.html#setting-the-right-airflow-user"
    echo
    export AIRFLOW_UID=$(id -u)
fi
one_meg=1048576
mem_available=$((($(getconf _PHYS_PAGES) * $(getconf PAGE_SIZE) / one_meg))
cpus_available=$(grep -cE 'cpu[0-9]+' /proc/stat)
disk_available=$(df / | tail -1 | awk '{print $4}')
warning_resources="false"
if (( mem_available < 4000 )) ; then
    echo
    echo -e "\033[1;33mWARNING!!!: Not enough memory available for Docker.\e[0m"
    echo "At least 4GB of memory required. You have $(numfmt --to iec $((mem_available * one_meg)))"
    echo
    warning_resources="true"
fi
if (( cpus_available < 2 )); then
    echo
    echo -e "\033[1;33mWARNING!!!: Not enough CPUS available for Docker.\e[0m"
    echo "At least 2 CPUS recommended. You have ${cpus_available}"
    echo
    warning_resources="true"

```

```

fi
if (( disk_available < one_meg * 10 )); then
    echo
    echo -e "\033[1;33mWARNING!!!: Not enough Disk space available for Docker.\e[0m"
    echo "At least 10 GBs recommended. You have $(numfmt --to iec $((disk_available * 1024 )))"
    echo
    warning_resources="true"
fi
if [[ ${warning_resources} == "true" ]]; then
    echo
    echo -e "\033[1;33mWARNING!!!: You have not enough resources to run Airflow (see above)!\e[0m"
    echo "Please follow the instructions to increase amount of resources available:"
    echo "    https://airflow.apache.org/docs/apache-airflow/stable/howto/docker-compose/index.html#before-you-begin"
    echo
fi
echo
echo "Creating missing opt dirs if missing:"
echo
mkdir -v -p /opt/airflow/{logs,dags,plugins,config}
echo
echo "Airflow version:"
/entrypoint airflow version
echo
echo "Files in shared volumes:"
echo
ls -la /opt/airflow/{logs,dags,plugins,config}
echo

```

```

    echo "Running airflow config list to create default config file if missing."
    echo
    /entrypoint airflow config list >/dev/null
    echo
    echo "Files in shared volumes:"
    echo
    ls -la /opt/airflow/{logs,dags,plugins,config}
    echo
    echo "Change ownership of files in /opt/airflow to ${AIRFLOW_UID}:0"
    echo
    chown -R "${AIRFLOW_UID}:0" /opt/airflow/
    echo
    echo "Change ownership of files in shared volumes to ${AIRFLOW_UID}:0"
    echo
    chown -v -R "${AIRFLOW_UID}:0" /opt/airflow/{logs,dags,plugins,config}
    echo
    echo "Files in shared volumes:"
    echo
    ls -la /opt/airflow/{logs,dags,plugins,config}

# yamllint enable rule:line-length
environment:
  <<: *airflow-common-env
  _AIRFLOW_DB_MIGRATE: 'true'
  _AIRFLOW_WWW_USER_CREATE: 'true'
  _AIRFLOW_WWW_USER_USERNAME: ${_AIRFLOW_WWW_USER_USERNAME:-airflow}
  _AIRFLOW_WWW_USER_PASSWORD: ${_AIRFLOW_WWW_USER_PASSWORD:-airflow}
  PIP_ADDITIONAL_REQUIREMENTS: ''

```

```

user: "0:0"

airflow-cli:
  <<: *airflow-common
  profiles:
    - debug
  environment:
    <<: *airflow-common-env
    CONNECTION_CHECK_MAX_COUNT: "0"
    # Workaround for entrypoint issue. See: https://github.com/apache/airflow/issues/16252
  command:
    - bash
    - -c
    - airflow
  depends_on:
    <<: *airflow-common-depends-on

# You can enable flower by adding "--profile flower" option e.g. docker-compose --profile flower up
# or by explicitly targeted on the command line e.g. docker-compose up flower.
# See: https://docs.docker.com/compose/profiles/
flower:
  <<: *airflow-common
  command: celery flower
  profiles:
    - flower
  ports:
    - "5555:5555"
  healthcheck:
    test: ["CMD", "curl", "--fail", "http://localhost:5555/"]

```

```

    interval: 30s
    timeout: 10s
    retries: 5
    start_period: 30s
    restart: always
    depends_on:
      <<: *airflow-common-depends-on
      airflow-init:
        condition: service_completed_successfully

volumes:
  postgres-db-volume:

```

3) (One-time) Initialize Airflow

This creates the database schema, default roles, and the admin user if configured:

```
docker-compose up airflow-init
```

You should see the `airflow-init` container **exit successfully**.

(If you forget this step, webserver logs complain: “You need to initialize the database. Please run `airflow db init`.”)

4) Start the Airflow stack

```
docker-compose up -d
```

```

PS C:\Airflow> docker-compose up -d
time="2025-08-18T17:00:01+05:30" level=warning msg="Found orphan containers ([airflow-airflow-worker-run-089ec51ff677 airflow-airflow-webserver-run-802484b8b097 airflow-airflow-a
pserver-1]) for this project. If you removed or renamed this service in your compose file, you can run this command with the --remove-orphans flag to clean it up."
[+] Running 8/8
 ✓ Container airflow-redis-1           Healthy         43.7s
 ✓ Container airflow-postgres-1        Healthy         43.7s
 ✓ Container airflow-airflow-triggerer-1 Running          0.0s
 ✓ Container airflow-airflow-webserver-1 Healthy         43.7s
 ✓ Container airflow-airflow-worker-1   Running          0.0s
 ✓ Container airflow-airflow-scheduler-1 Running          0.0s
 ✓ Container airflow-airflow-init-1     Exited          42.5s
 ✓ Container airflow-airflow-dag-processor-1 Started          0.0s
oint.s..." redis                About an hour ago   Up 19 minutes (healthy)    6379/tcp

```

Check status:

```
docker-compose ps
```



```
PS C:\Airflow> docker-compose ps
```

NAME	IMAGE	COMMAND	SERVICE	CREATED	STATUS	PORTS
airflow-airflow-dag-processor-1	apache/airflow:2.4.2	"/usr/bin/dumb-init ..."	airflow-dag-processor	2 minutes ago	Restarting (1) 1 second ago	
airflow-airflow-scheduler-1	apache/airflow:2.4.2	"/usr/bin/dumb-init ..."	airflow-scheduler	2 minutes ago	Up About a minute (healthy)	8080/tcp
airflow-airflow-triggerer-1	apache/airflow:2.4.2	"/usr/bin/dumb-init ..."	airflow-triggerer	2 minutes ago	Up About a minute (healthy)	8080/tcp
airflow-airflow-webserver-1	apache/airflow:2.4.2	"/usr/bin/dumb-init ..."	airflow-webserver	2 minutes ago	Up About a minute (healthy)	0.0.0.0:8080->8080/tc
airflow-airflow-worker-1	apache/airflow:2.4.2	"/usr/bin/dumb-init ..."	airflow-worker	2 minutes ago	Up 24 seconds (health: starting)	8080/tcp
airflow-postgres-1	postgres:13	"docker-entrypoint.s..."	postgres	2 minutes ago	Up 2 minutes (healthy)	5432/tcp
airflow-redis-1	redis:7.2-bookworm	"docker-entrypoint.s..."	redis	2 minutes ago	Up 2 minutes (healthy)	6379/tcp

You should see these services **running/healthy**:

- **airflow-webserver**
- **airflow-scheduler**
- **airflow-worker**
- **airflow-triggerer**
- **postgres**
- **redis**
- (**airflow-dag-processor** may restart; it's optional and not critical.)

5) Create an Admin user (PowerShell syntax!)

If you didn't set admin creds via env vars, create one **on a single line** (no \ in PowerShell):

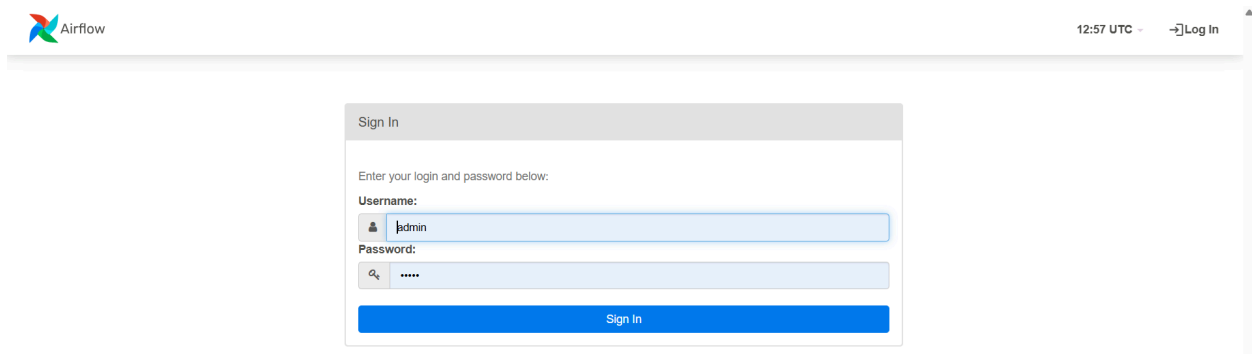
```
docker-compose run --rm airflow-worker airflow users create --role Admin --username admin --email admin --firstname admin --lastname admin --password admin
```

```
PS C:\Airflow> docker-compose run --rm airflow-worker airflow users create --role Admin --username admin --email admin --firstname admin --lastname admin --password admin
time="2025-08-18T17:02:36+05:30" level=warning msg="Found orphan containers ([airflow-airflow-worker-run-089ec51ff677 airflow-airflow-webserver-run-802484b8b097 airflow-airflow-a
pserver-1]) for this project. If you removed or renamed this service in your compose file, you can run this command with the --remove-orphans flag to clean it up."
[+] Creating 3/3
✓ Container airflow-postgres-1      Running      0.0s
✓ Container airflow-redis-1        Running      0.0s
✓ Container airflow-airflow-webserver-1 Running      0.0s
[+] Running 3/3
✓ Container airflow-postgres-1      Healthy      2.0s
✓ Container airflow-redis-1        Healthy      2.0s
✓ Container airflow-airflow-init-1  Exited      79.0s
```

Note: The multiline example with `\` is for bash. In PowerShell, use one line (or PowerShell's backtick ``` for line continuation).

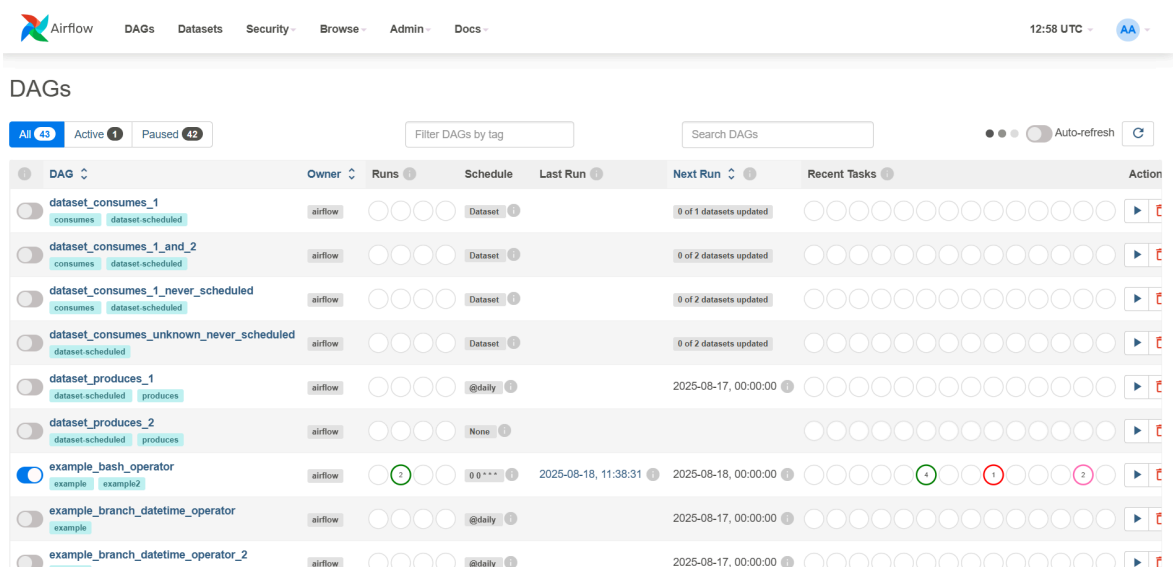
6) Open the Airflow UI

- Go to: <http://localhost:8080>
- Log in with: **admin / admin** (or whatever you created).



7) Verify with a test DAG

- In the UI, unpause any example DAG (if enabled) or drop a simple DAG file into `C:\Airflow\dags`.
- Trigger it and view task logs to confirm the worker is running tasks.



Troubleshooting we hit (and fixes)

1. argument GROUP_OR_COMMAND: invalid choice: 'api-server'

- Cause: Compose file referenced **airflow-apiserver** with command **api-server** (only in newer Airflow).
- Fix: **Remove** the **airflow-apiserver** service and **remove** any **depends_on: airflow-apiserver**.

2. line N: mapping key "healthcheck" already defined

- Cause: Duplicate **healthcheck** key in the same YAML block.
- Fix: Keep only one **healthcheck** section per service.

3. Web UI ERR_EMPTY_RESPONSE / webserver unhealthy

- Cause: Database wasn't initialized yet.
- Fix: Run **docker-compose up airflow-init**, then **docker-compose up -d**.

4. PowerShell parsing errors on --role etc.

- Cause: Using bash line continuations (****) in PowerShell.
- Fix: Run the whole command on **one line**.

5. Orphan containers warning

- Cause: Old containers from a previous compose definition.
- Fix:

docker-compose down --remove-orphans

docker-compose up -d