

Coding Assessment

Building a Simple Data Pipeline with Apache Airflow

Introduction

Apache Airflow is a workflow orchestration tool designed to automate, schedule, and monitor complex data processes. Instead of manually running scripts, Airflow lets you organize them into workflows called **DAGs (Directed Acyclic Graphs)**. It is widely used in data engineering for ETL pipelines, analytics, and machine learning tasks.

Core Features of Airflow

1. DAG-based workflows

- Workflows are represented as DAGs, where tasks are connected in a clear order without loops.

2. Powerful Scheduling

- You can schedule jobs using CRON expressions, intervals, or manual triggers.
- Missed jobs can be backfilled easily.

3. Extensible Operators

- Airflow provides operators like `PythonOperator`, `BashOperator`, and SQL operators to cover most data tasks.
- You can also write custom operators.

4. Scalability

- Supports distributed execution with Celery, Kubernetes, or Local executors.

5. User Interface

- A clean web UI to track DAGs, view logs, retry tasks, and monitor system health.

6. Integration Support

- Works with databases, cloud platforms, data lakes, APIs, and more.

Steps to Build a Pipeline

Step 1: Environment Setup

We'll use Docker Compose to launch Airflow locally.

Get the official docker-compose file

<https://airflow.apache.org/docs/apache-airflow/stable/docker-compose.yaml>

Create required directories

```
mkdir -p ./dags ./logs ./plugins
```

Set user environment variable

```
echo -e "AIRFLOW_UID=$(id -u)" > .env
```

Initialize database

```
docker compose up airflow-init
```

```
PS C:\Airflow> docker compose up airflow-init
[+] Running 2/2
 ✓ Container airflow-postgres-1  Running      0.0s
 ✓ Container airflow-redis-1     R...      0.0s
Attaching to airflow-init-1
airflow-init-1 |
airflow-init-1 | WARNING!!!: Not enough memory available for Docker.
airflow-init-1 | At least 4GB of memory required. You have 3.8G
airflow-init-1 |
airflow-init-1 | WARNING!!!: You have not enough resources to run Airflow (see above)!
airflow-init-1 | Please follow the instructions to increase amount of resources available:
airflow-init-1 | https://airflow.apache.org/docs/apache-airflow/stable/howto/docker-compose/index.html#before-you-begin
airflow-init-1 |
airflow-init-1 | Creating missing opt dirs if missing:
airflow-init-1 |
airflow-init-1 | Airflow version:
airflow-init-1 | The container is run as root user. For security, consider using a regular user account.
airflow-init-1 |
```

Start all services

```
docker compose up
```

```

PS C:\airflow> docker compose up
>>
time="2025-08-19T12:20:42+05:30" level=warning msg="Found orphan containers ([airflow-airflow-worker-run-7d5863164f41]) for this project. If you removed or renamed this service in your compose file, you can run this command with the --remove-orphans flag to clean it up."
[+] Running 7/7
 ✓ Container airflow-redis-1           Running      0.0s
 ✓ Container airflow-postgres-1        Running      0.0s
 ✓ Container airflow-airflow-triggerer-1 Running      0.0s
 ✓ Container airflow-airflow-scheduler-1 Running      0.0s
 ✓ Container airflow-airflow-webserver-1 Running      0.0s
 ✓ Container airflow-airflow-dag-processor-1 Running      0.0s
 ✓ Container airflow-airflow-worker-1  Running      0.0s
Attaching to airflow-dag-processor-1, airflow-init-1, airflow-scheduler-1, airflow-triggerer-1, airflow-webserver-1, airflow-worker-1, postgres-1, redis-1

```

Open the UI at <http://localhost:8080> and log in:

- Username: **airflow**
- Password: **airflow**

Step 2: Create a Postgres Connection

In the UI, navigate to Admin → Connections and add:

- **Conn ID:** tutorial_pg_conn
- **Type:** Postgres
- **Host:** postgres
- **Database:** airflow
- **Login:** airflow
- **Password:** airflow
- **Port:** 5432

Edit Connection

Connection Id *	tutorial_pg_conn
Connection Type *	Postgres
	Connection Type missing? Make sure you've installed the corresponding Airflow Provider Package.
Description	
Host	postgres
Database	airflow
Login	airflow
Password	*****
Port	5432

Step 3: Define Tables

We'll create two tables:

- `employees_raw` → temporary staging table
- `employees_clean` → final processed table

These will be created using the `SQLExecuteQueryOperator`.

Step 4: Download and Load Data

Use Python tasks to download a CSV file and insert it into the staging table. Airflow hooks handle the database interaction.

Step 5: Merge and Clean Data

A second task will merge deduplicated data into the final table using `INSERT ... ON CONFLICT`.

Example DAG

```
import datetime

import pendulum

import os

import requests


from airflow.decorators import dag, task

from airflow.providers.postgres.hooks.postgres import PostgresHook

from airflow.providers.postgres.operators.postgres import PostgresOperator


@dag(

    dag_id="process_employees",
```

```


        schedule="0 0 * * *",

        start_date=pendulum.datetime(2021, 1, 1, tz="UTC"),

        catchup=False,

        dagrun_timeout=datetime.timedelta(minutes=60),
    )

def ProcessEmployees():

    #  Create employees table

    create_employees_table = PostgresOperator(

        task_id="create_employees_table",

        postgres_conn_id="tutorial_pg_conn",

        sql="""

            CREATE TABLE IF NOT EXISTS employees (

                "Serial Number" NUMERIC PRIMARY KEY,

                "Company Name" TEXT,

                "Employee Markme" TEXT,

                "Description" TEXT,

                "Leave" INTEGER

            );

        """,

    )

    create_employees_temp_table = PostgresOperator(

        task_id="create_employees_temp_table",

```

```

postgres_conn_id="tutorial_pg_conn",

sql="""

    DROP TABLE IF EXISTS employees_temp;

    CREATE TABLE employees_temp (

        "Serial Number" NUMERIC PRIMARY KEY,

        "Company Name" TEXT,

        "Employee Markme" TEXT,

        "Description" TEXT,

        "Leave" INTEGER

    );

    """,

)

@task

def get_data():

    data_path = "/opt/airflow/dags/files/employees.csv"

    os.makedirs(os.path.dirname(data_path), exist_ok=True)

    url =
"https://raw.githubusercontent.com/apache/airflow/main/airflow-core/docs/tutorial/pipeline_example.csv"

    response = requests.get(url)

    with open(data_path, "w") as file:

```

```

        file.write(response.text)

    postgres_hook = PostgresHook(postgres_conn_id="tutorial_pg_conn")

    conn = postgres_hook.get_conn()

    cur = conn.cursor()

    with open(data_path, "r") as file:

        cur.copy_expert(

            "COPY employees_temp FROM STDIN WITH CSV HEADER DELIMITER
AS ',' QUOTE '\"',

            file,

        )

    conn.commit()

@task

def merge_data():

    query = """

        INSERT INTO employees

        SELECT *

        FROM (

            SELECT DISTINCT *

            FROM employees_temp

        ) t

        ON CONFLICT ("Serial Number") DO UPDATE

        SET

```

```

        "Employee Markme" = excluded."Employee Markme",

        "Description" = excluded."Description",

        "Leave" = excluded."Leave";

    """

    try:

        postgres_hook =
PostgresHook(postgres_conn_id="tutorial_pg_conn")

        conn = postgres_hook.get_conn()

        cur = conn.cursor()

        cur.execute(query)

        conn.commit()

        return 0

    except Exception as e:

        return 1

# Task dependencies

[create_employees_table, create_employees_temp_table] >> get_data() >>
merge_data()

dag = ProcessEmployees()

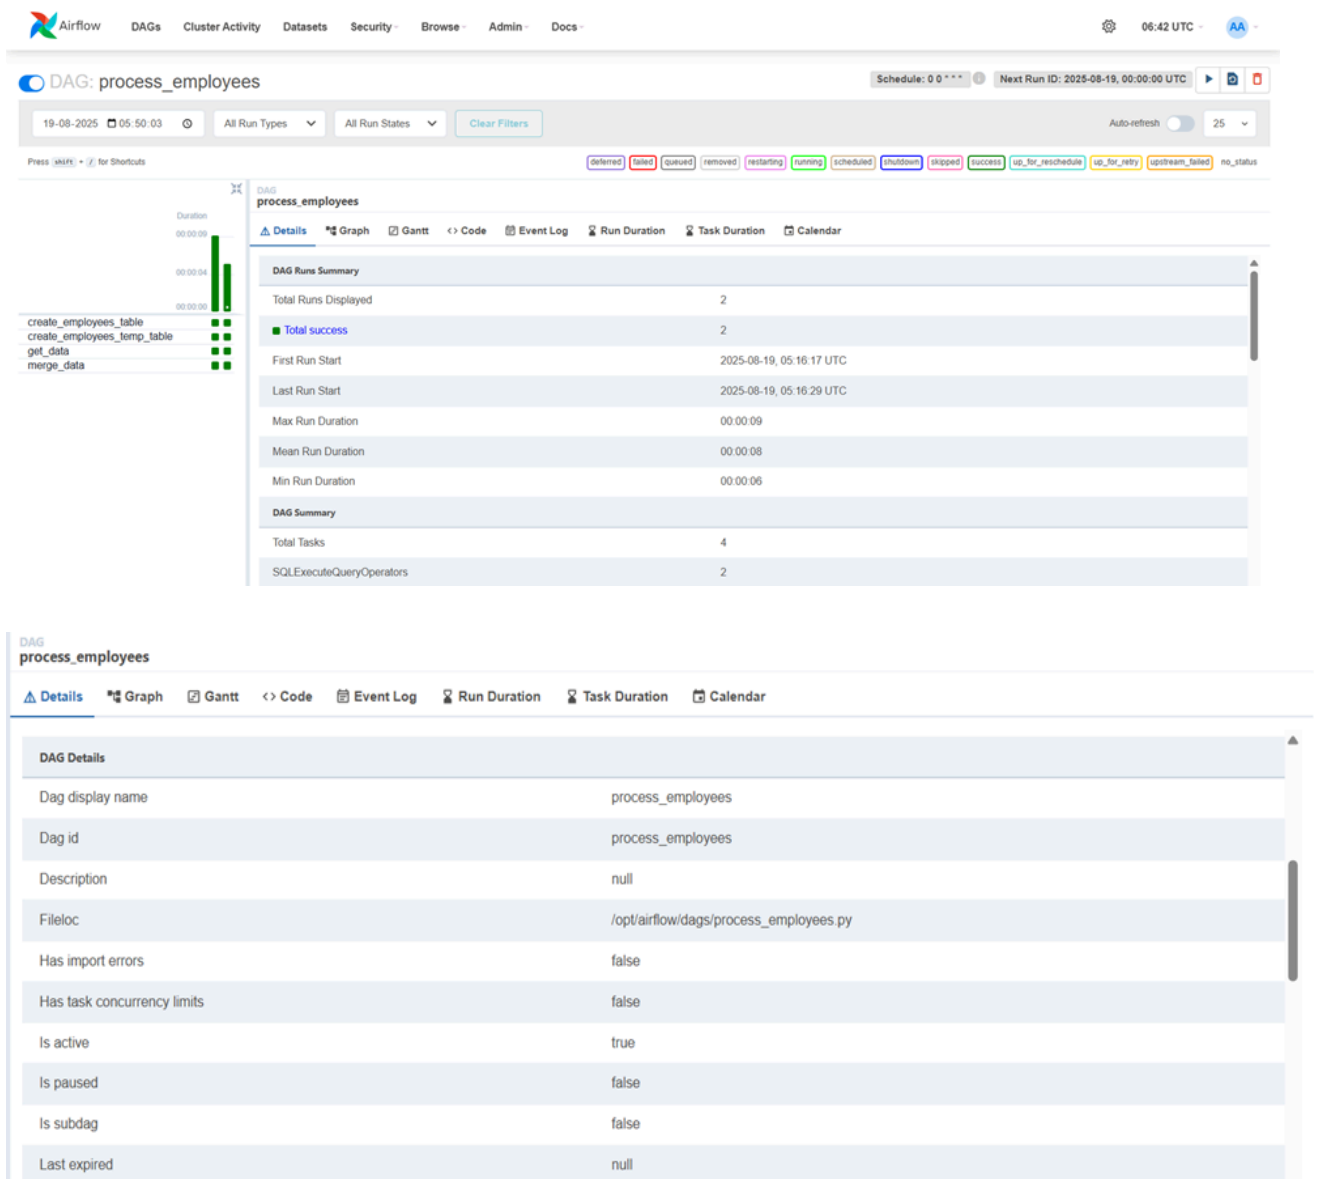
```


Running the DAG

Once the DAG was saved, it appeared in the Airflow UI. The DAG was triggered manually, and the pipeline successfully:

- Downloaded the CSV file
- Inserted the data into the staging table
- Merged and cleaned the data into the final table

This confirmed that the data pipeline was functioning correctly end-to-end.



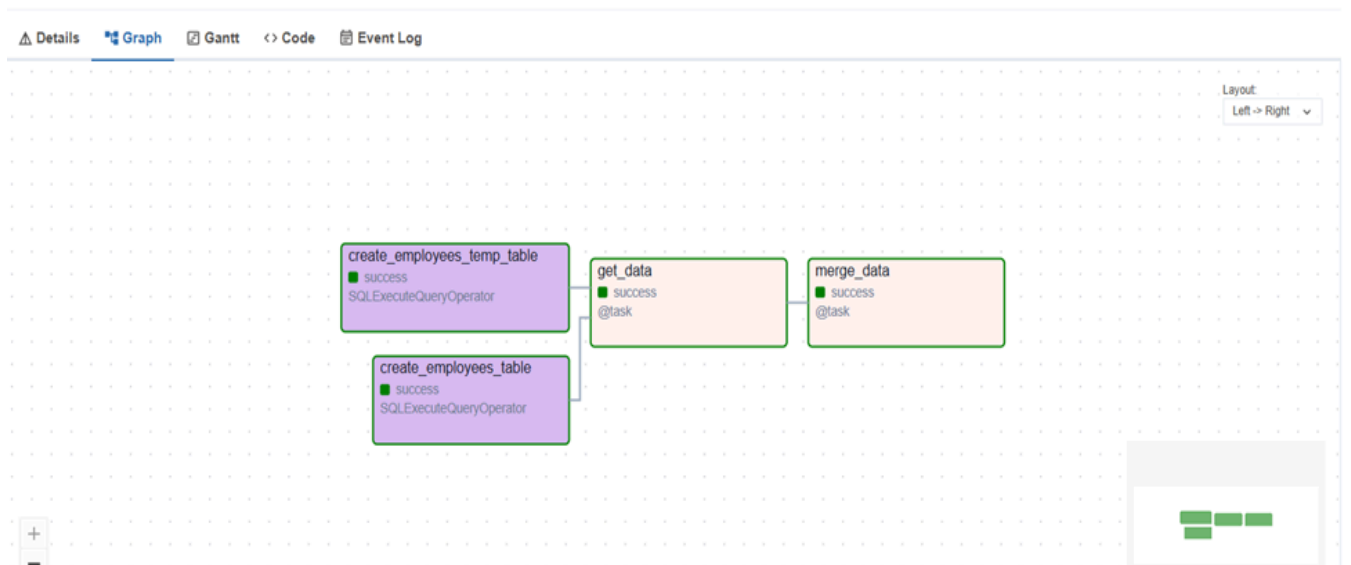
DAG
process_employees

[Details](#)
[Graph](#)
[Gantt](#)
[Code](#)
[Event Log](#)
[Run Duration](#)
[Task Duration](#)
[Calendar](#)

End date	null
Is paused upon creation	null
Last parsed	2025-08-19T05:43:06.000644+00:00
Orientation	LR
Render template as native obj	false
Start date	2021-01-01T00:00:00+00:00
Template searchpath	null
Timezone	UTC
Owners	
Tags	No tags
Schedule interval	0 0 * * *
Dag run timeout	3600 seconds

Exploring Airflow Views

- **Graph View** → Displays task dependencies in a node graph format.
- **Gantt View** → Timeline view of task duration and overlaps.
- **Tree View** → Quick overview of success/failure across historical runs.




DAG Run
process_employees / 2025-08-19, 00:00:00 UTC

Clear

Mark state as...

Details Graph **Gantt** Code Event Log



localhost:8080/admin/airflow/tree?dag_id=tutorial

Google Lens

Airflow DAGs Data Profiling Browse Admin Docs About

2025-08-20 04:32:14 UTC

DAG: tutorial

schedule: 1 day, 0:00:00

Graph View **Tree View** Task Duration Task Tries Landing Times Gantt Details Code Trigger DAG Refresh Delete

Base date: Number of runs: 25 Go

BashOperator

success running failed skipped upstream_failed up_for_reschedule up_for_retry queued no_status

